

# **Inferring Interestingness in Online Social Networks**

**Will Webberley**

**2013**

**Cardiff University  
School of Computer Science & Informatics**

**Declaration**

This work has not previously been accepted in substance for any degree and is not concurrently submitted in candidature for any degree.

Signed ..... (candidate)

Date .....

**Statement 1**

This thesis is being submitted in partial fulfillment of the requirements for the degree of PhD.

Signed ..... (candidate)

Date .....

**Statement 2**

This thesis is the result of my own independent work/investigation, except where otherwise stated. Other sources are acknowledged by explicit references.

Signed ..... (candidate)

Date .....

**Statement 3**

I hereby give consent for my thesis, if accepted, to be available for photocopying and for inter-library loan, and for the title and summary to be made available to outside organisations.

Signed ..... (candidate)

Date .....

Copyright © 2013 Will Webberley.

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.2 or any later version published by the Free Software Foundation; with no Invariant Sections, no Front-Cover Texts, and no Back-Cover Texts

**To People you care  
for their patience and support.**

# Abstract

Abstract text here etc.

# Contents

<b>Abstract</b>	<b>iv</b>
<b>Contents</b>	<b>v</b>
<b>List of Publications</b>	<b>x</b>
<b>List of Figures</b>	<b>xi</b>
<b>List of Tables</b>	<b>xiii</b>
<b>List of Acronyms</b>	<b>xiv</b>
<b>Glossary</b>	<b>xv</b>
<b>1 Introduction</b>	<b>1</b>
1.1 Twitter and Retweeting . . . . .	1
1.2 The Problem . . . . .	2
1.3 Contributions . . . . .	2
1.4 Thesis Structure . . . . .	2

---

<b>2</b>	<b>Background</b>	<b>3</b>
2.1	Domain Context . . . . .	3
2.2	Research Motivation . . . . .	4
<b>3</b>	<b>Related Literature</b>	<b>5</b>
3.1	Twitter and Retweeting . . . . .	5
3.2	Information Quality and Retrieval . . . . .	5
3.3	Tweet Interestingness . . . . .	6
3.4	Scoring Tweets . . . . .	7
<b>4</b>	<b>Understanding The Behaviour of Retweeting in Twitter</b>	<b>8</b>
4.1	Retweet Trees . . . . .	10
4.2	Information Retrieval . . . . .	11
4.2.1	Epistemic Search . . . . .	12
4.2.2	Hedonic Search . . . . .	12
4.2.3	Affective Stimulation . . . . .	12
4.2.4	Search ‘bubble’ . . . . .	12
4.3	Background to This Area . . . . .	12
4.4	Twitter Analysis . . . . .	14
4.5	Experimental Methods . . . . .	15
4.5.1	Data Collection Methodology . . . . .	15
4.5.2	Data . . . . .	16
4.6	Empirical Results . . . . .	17

---

4.6.1	Length of Retweet Chains . . . . .	17
4.6.2	Size of Retweet Groups . . . . .	19
4.6.3	Retweet Follower Pattern . . . . .	22
4.6.4	Retweet Time Delay . . . . .	26
4.7	Summary . . . . .	27
<b>5</b>	<b>Analysis of Twitter's Social Structure</b>	<b>29</b>
5.1	Background to This Area . . . . .	31
5.2	Overview . . . . .	31
5.3	Model . . . . .	32
5.3.1	Machine Learning . . . . .	32
5.3.2	Algorithm . . . . .	34
5.3.3	Features . . . . .	35
5.4	Training the Model . . . . .	36
5.4.1	Data Collection . . . . .	36
5.4.2	Feature Extraction and Regression Training . . . . .	36
5.5	Network Analyses . . . . .	36
5.5.1	Path Network . . . . .	36
5.5.2	Random Network . . . . .	38
5.5.3	Scale-Free Network . . . . .	38
5.5.4	Comparison to Real Twitter Data . . . . .	39
5.5.5	Structure Comparison . . . . .	39



---

5.6	Interesting Additional Findings . . . . .	41
5.6.1	Graph Density . . . . .	41
5.6.2	Results . . . . .	41
5.6.3	Uses . . . . .	41
5.7	Predictions From the User Graph . . . . .	42
5.7.1	Data Collection . . . . .	42
5.7.2	Validating Results . . . . .	44
5.7.3	Improving This . . . . .	45
5.8	Future Work . . . . .	46
5.9	Summary . . . . .	46
<b>6</b>	<b>Inferring Interestingness of Tweets based on Information Flow Through the Network</b>	<b>48</b>
6.0.1	User Influence . . . . .	51
6.1	Interestingness Scores . . . . .	51
6.1.1	‘Twitter is a Memepool’ . . . . .	53
6.2	Retweet Volumes as Nominal Attributes . . . . .	55
6.2.1	Binning the Retweet Outcomes . . . . .	56
6.2.2	Varying Bin Sizes . . . . .	57
6.3	Techniques Used . . . . .	57
6.3.1	Machine Learning . . . . .	57
6.3.2	Crowdsourcing . . . . .	58
6.4	Experimentation . . . . .	58

---

6.4.1	Data Collection . . . . .	58
6.4.2	Data Corpora . . . . .	59
6.4.3	Features . . . . .	60
6.4.4	Making predictions and inferring Tweet interestingness . . . .	62
6.4.5	Planning the Result Validation . . . . .	62
6.4.6	Validation Results . . . . .	64
6.5	Further Analyses . . . . .	67
6.5.1	Discerning Interesting Information from Noise . . . . .	67
<b>7</b>	<b>Critical Assessment and Conclusions</b>	<b>70</b>
7.1	Critical Analysis of Results . . . . .	70
7.1.1	Analysis of Initial Research . . . . .	70
7.1.2	Analysis of Final Results . . . . .	70
7.2	Further and Future Work . . . . .	71
7.3	Conclusions . . . . .	71
7.3.1	Summary . . . . .	71
7.3.2	Contributions . . . . .	71
	<b>Appendices</b>	<b>72</b>
	<b>Bibliography</b>	<b>73</b>

## List of Publications

Some of the work produced towards this thesis has also been published separately as follows.

- W. Webberley, S. M. Allen, R. M. Whitaker. *Inferring the Interesting Tweets in Your Network*, in *Workshop on Analyzing Social Media for the Benefit of Society (SOCIETY 2.0)*, 3<sup>rd</sup> International Conference on Social Computing and its Applications (SCA), Karlsruhe, Germany. *IEEE 2013*
- W. Webberley, S. Allen, R. Whitaker. *Retweeting: A Study of Message-Forwarding in Twitter*, in *Workshop on Mobile and Online Social Networks (MOSN'11)*, 5<sup>th</sup> International Conference on Network and System Security (NSS), Milan, Italy. *IEEE 2011*

## List of Figures

4.1	<i>Log/log distribution of maximum path-lengths from retweet groups. . .</i>	17
4.2	<i>Maximum likelihood power-law fit for the cumulative distribution of retweet group sizes. . . . .</i>	20
4.3	<i>Log/log relationship between the maximum path-length and size of a retweet group. . . . .</i>	20
4.4	<i>Relationship between a retweet group's audience size and its longest path-length. . . . .</i>	21
4.5	<i>Relationship between number of followers (and respective distribution) of the original tweeter as the path-length increases. . . . .</i>	24
4.6	<i>Proportion of final retweeters following upstream users at varying distances along the chain. . . . .</i>	25
4.7	<i>Average time (in seconds) between first post and final retweet of a retweet group varying with the group's maximum path-length. . . . .</i>	26
5.1	Retweet volume frequency distribution from path network simulation	37
5.2	Retweet volume frequency distribution from random network simulation	38
5.3	Comparing the retweet volumes distribution from scale-free graph simulation to data from Twitter's graph . . . . .	40

6.1	Proportionate frequency distribution of global scores in the entire test set compared to only those in the subset $T$ . . . . .	64
6.2	The probability of the MTW's chosen Tweet's $TS_{Avg}(t)$ being in the top $n$ of Tweets for that question while varying the minimum allowed maximum Tweet score of the question. . . . .	65
6.3	The probability that Tweet $t$ is chosen given that $TS_G(t)$ is greater than a given value, $x$ . . . . .	66
6.4	The probability of a confident selection being made for question $q$ with varying $d(q)$ . . . . .	69

## List of Tables

6.1	Features used to train the model from the global data corpus . . . . .	61
6.2	Absolute $TS_{Avg}(t)$ disparity of questions with varying number of confident answers made. Entries in <b>bold</b> are used to highlight interesting values. . . . .	67
6.3	Score disparity comparison between selected Tweets of question $q$ and <i>all</i> Tweets in $q$ when using the three different Tweet scores as metrics	69

# List of Acronyms

**OSN** Online Social Network

**MTW** Mechanical Turk Worker

# Glossary

**Follower** A follower is a type of user. A user,  $x$ , is a follower of user  $y$  if user  $x$  follows user  $y$ . Other users who follow a particular user will receive all of the user's Tweets and retweets to their home timeline. A user can elect to follow another user.

**Friend** The inverse of follower. A friend is a user that another user follows.

**Path-length** The penetration of a Tweet - i.e. the number of times a Tweet is retweeted down one chain. The final retweeter in the chain indicates the number of hops the Tweet has taken from its author.

**Retweet**  $n$ . - A replica of a Tweet, which has been forwarded on by a user (who is not the Tweet's original author) to their own followers.

$v$ . - The act of replicating a Tweet. A user who finds a Tweet interesting may retweet it so that it gets more exposure.

**Retweet Group** Set of Twitter users responsible for the propagation of a Tweet. Comprises the original author of the Tweet and the users which have retweeted it.

**Retweet Volume** The number of times a particular Tweet has been retweeted.

**Timeline** A collection of Tweets in Twitter in reverse-chronological order. A user's timeline consists of that user's Tweets. A user's home timeline consists of the Tweets of each friend of the user.

**Tweet**  $n$ . - A piece of information in Twitter; a piece of text, less than 240 characters long, which is written by a user. When sent, the Tweet is sent to the home timelines of



each of the followers of the Tweet's author.

v. - The act of writing and sending a Tweet.

**User** An account on Twitter. Each user (usually representing a real-life person or organisation) can Tweet, retweet, follow other users and be followed by other users.

---

# Chapter 1

## Introduction

Online social networks have exploded into the lives of millions of people worldwide over the last decade, and their use has dominated the communication highways and facilitated the interconnection of the world in ways never before perceived possible.

These social networks imitate real-world social networks. Although most such platforms each provide a different service to collaboratively satisfy an array of different use-cases, they tend to all be based around the idea of ‘friendships’ (i.e. links between the user nodes in the social graph) and the sharing of information amongst friends.

Social networks like these have been available for around ten years now (with MySpace<sup>1</sup> launching in 2003 and Bebo<sup>2</sup> in 2005), but it wasn’t really until Facebook’s<sup>3</sup> world-wide launch in 2006 that social networks became the staple, ubiquitous norm that they are today. More recently, we have seen the introductions of Google’s social network grown from its Buzz service, Google Plus, Pinterest, App.net, and many more. They make up a large part of the basis and meaning behind the ideas of Web 2.0, which describes the web as being primarily formed from user-generated content and encourages the sharing of such content.

---

<sup>1</sup><http://myspace.com>

<sup>2</sup><http://bebo.com>

<sup>3</sup><http://facebook.com>

## 1.1 Twitter as a Social Network

Twitter is an online social network, which launched in the summer of 2006. Since then, it has rapidly gained in popularity amongst several different user groups - teens and young people, casual users, celebrities, reporters, and so on. Although Twitter has never been a direct competitor with Facebook, users tend to use the two sites concurrently for different purposes: whilst Facebook's focus is on providing many services at once (such as photo-sharing, commenting/endorsing of information, messaging, pages for businesses, groups, events, etc.), Twitter's is more on simplicity.

Further to being an online social network, Twitter is a microblogging website. Whilst a blog ('web-log') typically contains posts containing long pieces of text or other media focusing on a single or set of topics, such as a diary, Twitter only allows its users to post short pieces of text, up to 140 characters in length, called Tweets. Whilst each Tweet may only realistically be able to hold a couple of sentences, this system facilitates quick, timely, and *live* information-sharing amongst its millions of users. Its idea is that short pieces of news will 'travel' faster and will be seen by more people more quickly than traditional news stories.

Although Tweets are limited to 140 characters in length, the inclusion of URLs is allowed. This enables further extension of Tweets through external websites, and supports the inclusion of links to images and videos. Twitter has encouraged this use-case by providing 'share' buttons for developers to embed in websites, and direct support for photo and video applications, such as TwitPic and Vine. Its simplicity has also helped its growth into the mobile domain, in which smartphone users are able to very quickly post updates about their lives, a piece information they want to share, or a photo or video, and be able to post it *as it happens*. This has been especially useful in emergency situations worldwide, including the Haiti earthquake in 2010 and the protests in Egypt in 2011.

## 1.2 Twitter's Social Graph

As with many social networks, the structure of Twitter lies within the users and their connectivity within its social graph. However, unlike Facebook, Twitter's primary social graph is made up of *directional* links between its users. A person using Twitter can elect to *follow* another user

## 1.3 Information Propagation through Retweeting

Introduction to Twitter, act of Tweeting and retweeting information. Ideas:

- Twitter is an online social network
- Microblogging - short, quick updates - useful for growing mobile domain
- links in tweets allow extension
- how the social graph is built up
- Used for quick information consumption in spare moments
- endorsement of info - retweets / likes / etc.
- Retweeing act of forwarding on of information to those who follow you

## 1.4 The Problem

Explain about the problem of 'noise' in Twitter:

- Not all information is interesting
- Noisy tweets can dampen experience - becomes harder to find interesting information

- Users cannot find what they do not know exists (similar to Google search bubble)
- etc.

## 1.5 Contributions

- How is this work novel?
- What benefits does this research provide?
- Does the work solve The Problem?
- etc.

## 1.6 Thesis Structure

Break down the structure of thesis (i.e. refer to contents page, but also a general overview of the order of sections and what is discussed).

---

## Chapter 2

# Background

## 2.1 Domain Context

Go into more information regarding information propagation in Twitter and about the mechanics of Twitter. Ideas:

- Retweet volumes
- communities and how networks form
- ‘search bubbles’
- propagation
- social structure (and its effect on propagation)
- information quality, relevance, retrieval, filtration
- information interestingness
- categorise above into subsections (possibilities: the network and communities, information retrieval and relevance, information retrieval and bubbles, and retweeting as a form of propagation)

## 2.2 Research Motivation

Link to ‘The Problem’ in introduction chapter. Talk about wanting to allow people to be exposed to information that they are *likely* to find interesting based on the interestingness of the tweet, but without them having to search for the information or follow the users responsible for sourcing or forwarding the information.

Further work would be done on refining this based on a per-user basis (i.e. that user’s particular interests as a relevance metric for the interestingness of the Tweet).

---

## **Chapter 3**

### **Related Literature**

Either overall lit-review here, individual ones for each chapter, or both (i.e. generally overall and specific ones too).

#### **3.1 Twitter and Retweeting**

References for work that helped with the introduction and initial understanding of Twitter. Include papers on:

- Twitter communities and network
- Influential users ('hubs' and 'authorities')

#### **3.2 Information Quality and Retrieval**

References for work regarding quality and relevance of information and retrieval techniques

- Recognition heuristic
- Google search 'bubble'



- How Twitter handles information quality (Twitter noise, relevance) (evangelists and detractors - Bighona)
- Document retrieval and recommendation systems (Gavalas)

### 3.3 Tweet Interestingness

Retweeting has become the focus of much research, particularly in defining the interest or credibility levels of information [3]. The phenomenon can be seen as a way in which users can ‘vote’ for information in Twitter. If a user views a Tweet that he/she finds interesting, or expects his/her followers to, the Tweet can be retweeted so that the audience of the information is extended. It’s unlikely that a user will retweet something uninteresting or inane, since they realise that their followers are probably not going to find it very engaging. Thus a user’s friends (i.e. *followees*) become filters of information by only forwarding on interesting information to their followers.

In [11], the idea of Twitter communities is discussed, and how clusters of users with similar interests can build up over time in the social graph. This means that it’s likely that several friends and followers of a user have a similar set of interests to the user, and are therefore likely to find similar Tweets of interest.

[16] discusses an idea similar to that of predicting the interestingness of a Tweet, but focuses largely on predicting the users most likely to find the Tweet interesting enough to retweet. Similarly, [10] looks at the same issue but at the other way around by predicting the *type* of Tweets that are likely to be retweeted many times. Discussions on the retweet decision in relation to a user’s recognition of its features take place in [5] as well as conclusions about the effect of features such as the number of followers of a user or any pre-existing metadata on the interestingness of the Tweet.

A model for making predictions on retweet volume probability is introduced by [20] (and [14]), which forms the basis for some of the ideas behind this paper. Their work involves the use of a logistic regression, trained by an extensive set of features, to pre-

dict retweet outcomes for a Tweet. We try and simplify parts of their methodology for efficiency, so that interestingness inferences can be made quickly and efficiently as the retweets occur in real time.

[13] and [16] also use machine learning to train models to predict retweet outcomes as part of their work, with the former using a logistic regression and employing their prediction forecasts to infer *what* makes a Tweet interesting and the latter using a decision tree classifier to draw links between Tweets that are interesting to a user and that user's retweet likelihood on those Tweets.

### 3.4 Scoring Tweets

<https://sites.google.com/site/learningtweetvalue/home> - these guys came up with their own method of predicting retweet outcome (do not relate to interestingness per se, and instead look at the past history of user's tweet segments (such as certain features used, what hashtags contained, etc.)

---

## ***Chapter 4***

# **Understanding The Behaviour of Retweeting in Twitter**

Stuff to finish up in this section:

- Explain motivation for research in this particular area
- Use this motivation to explain the purpose for this research as a basis for the work in the next few chapters
- Explain how this chapter is the basis for research into Twitter's social structure in the next chapter
- Normalise terms (retweet-group size / retweet volume) here and in further chapters throughout thesis
- (e.g. we have addressed tweet quality in terms of propagation, can a network have a quality too? what further factors can affect the dissemination of information in social networks?)

Social networking and communication now make up a significant part of the Internet, with sites such as Facebook and MySpace attracting millions of users worldwide. Blogging websites, such as Tumblr and Wordpress, have also seen a significant influx of users, allowing people to update and share posts and links on their interests and

every day lives. However, as the need for quick and short updates increases, microblogging sites, such as Twitter, have become more and more popular.

Microblogging is a form of blogging in which posts are limited to a specific character count (140 in the case of Twitter). This limit means that users tend to post updates more frequently [19]. The furthering development of applications in the mobile domain, particularly for iOS, Android and Symbian systems, mean that users update from wherever they are and whenever they like.

Typically, users will tweet of topics that interest them. This may be related to their work, a hobby, or a mixture of multiple areas. These tweets are generally posted with the idea that they will be useful or interesting for some of their followers as well as an attempt to attract more followers. Zhao et al. [19], showed that Twitter is also used as a means to contact friends and to get assistance and opinions on topics. Therefore, particular users may belong to different communities of people depending on what kind of posts they want to view.

The strength of Twitter is in its social structure, where users can elect to follow others. Followers of a user receive all of that user's posts in their individual (or 'home') timelines. If a user has set their profile to be public, then their posts also appear on the public timeline, which is accessible to anyone; even those without a Twitter account. As a result, people are likely to follow users who update with interesting posts; whether the follower is a big fan of the friend (a 'friend' is said to be a *followee* of a user) and simply wants to know everything going on in their life, or if the follower is simply interested in the topical area of most of the friend's posts. If someone sees a post that they feel would be interesting to their followers, they can 'retweet' the post, which then relays it as a tweet.

The friends of a particular user effectively become *filters* of information for that user. The user can choose to follow another user, and therefore implicitly indicates the kind of information they want to receive. If you, a Twitter user, want to gain more attention by means of followers, then by posting interesting posts (or, at least, posts that others want to read), you will

1. increase the chances that users reading your posts will choose to follow you, and
2. increase the chances that users will decide to retweet your post, thus broadcasting your tweet to a larger audience. People viewing the retweet then may decide to follow you - the original tweeter.

It is, of course, possible that more than one user may decide to retweet one of your posts. A follower of a retweeter may decide to retweet the retweet, creating a chain. Naturally, the larger your effective audience (both directly and through retweets), the greater chance, again, you have of being retweeted. [15] showed this by demonstrating how the retweet rate increases with the number of followers of the original tweeter. This is also related to the ideas of user influence mentioned in [4]; that more followers implies more influence. The retweeting process ultimately results in a tree system, with the root being the original tweeter, and other nodes being retweeters at different levels, similar to the idea of URL cascades discussed in [8].

The motivation for this study has stemmed largely from the interesting decentralised nature of retweets, spurred on by similar literature such as information cascades in [8]. The notion of retweeting is also very similar to content forwarding in opportunistic networks, particularly ideas discussed in [1], and studying the mechanisms of retweeting from a human-centric perspective may provide insight into protocols for autonomous content dissemination. We hope that understanding retweet properties may help us in forwarding relevant information to a user from without that user's social neighbourhood.

## 4.1 Retweet Trees

Example of a retweet tree.

Source users are roots, final retweeters are leaves (although subject to change as time goes by).

Links between nodes indicate retweeting from a user to another user (but not which users follow which).

Define path-length and retweet groups.

For this paper, the term *path-length* is to be defined as the number of times a particular tweet is retweeted down a chain. For example, the status;

```
RT @user1: RT @user2: This is the body of the tweet
```

has a path length of two. All users involved in this retweet chain have used the old method of retweeting, which is still very commonly used and has effectively provided most of the results for this paper.

## 4.2 Information Retrieval

People don't use Twitter to obtain any specific information (they don't know what people are necessarily going to say)

If all friends are technology based, then you can expect technological tweets, but generally this isn't the case.

#### 4.2.1 Epistemic Search

#### 4.2.2 Hedonic Search

#### 4.2.3 Affective Stimulation

#### 4.2.4 Search ‘bubble’

### 4.3 Background to This Area

To date, research has largely focused on the actual structure of Twitter in terms of the friend-follower graph [?], user communities [11] and also the psychology behind Twitter users [19] [4]. A lot of these studies link well with retweeting, as will be explained later, and researching them has helped to explain and facilitate retweet studying at a wider angle.

Arvola et al. [2] introduced a retrieval system for online documents which returns particular passages in documents which are of a pre-defined interest for the user (generally through the use of search terms). In Twitter, a user can specify what they are interested in receiving by following a set of users, with their neighbours effectively becoming the search term. Twitter then enables the user to receive more relevant and interesting information based on these filters. However Twitter doesn’t alter the reading order of passages (or tweets) as in the aforementioned retrieval system, the home timeline only shows the tweets coming from people the user has chosen to follow. [18] discusses user’s relevance judgement based on hedonic and epistemic search, particularly linking to the idea of ‘affective stimulation’, occurring when the user becomes affected by some information (e.g. emotionally or for entertainment). This is somewhat also applicable to following people in Twitter as users don’t generally use Twitter to receive any particular form of information, thus they haven’t *searched* for any specific topic. Therefore, any information they do receive, and which they find interesting, be-

comes affective stimulation. If someone follows many users, then he/she is receiving too much information to effectively filter out the interesting information. If following a few users, he/she can't be receiving *all* the information that they might want to.

However, despite the fact that this system is thus quite sound (by largely displaying information the user *chose* to receive), it may not be particularly complete since it is likely that there are many other users and topics the user might still be interested in which is not returned by this Twitter information retrieval system. [2] also discusses the notion of expected reading effort and its importance in information retrieval systems. The character-precision recall metric was introduced as a way of demonstrating the tolerance-to-irrelevance ratio. The general mechanism for this is to do with users reading a document passage; the point the user reaches this ratio is when the user stops reading the particular passage and moves to the next whole document (since they assume the rest of the document is also irrelevant to them). Therefore, the more effective the information retrieval system, the lower the chance that this ratio is reached by the user. Associating this with Twitter, we can say that a user viewing tweets from a user they are following may get to the point where he or she reaches this ratio (i.e. is bored or finding the tweets irrelevant) and decides to unfollow the friend. Similar to above, the more effective the user is when selecting people to follow (friends) in the hope of receiving interesting information, the lower the chance that the user will remove these friends.

As discussed in the introduction, a user is more likely to follow other people who post tweets of interest to the user. As more and more people follow a particular user and start tweeting a lot about a particular topic, the users may start to follow each other, producing a growing 'swarm' of interest around a certain topic (or a particular user). As more users start associating with this user or topic, its audience becomes more widespread, attracting further users to it. As this continues over time, a community of users can be formed. This leads onto the notion of Twitter communities discussed and experimented on by [11], who define communities to be groups of users with dense follower links between them. Their studies continue into looking at the development



of communities over time and also defines how they can be formed from different types of people - such as those who tweet lots and have many followers, and those who post little and have few followers. Studies in the behaviour of different types of users in Twitter is done more thoroughly in [?], where ‘broadcasters’ - users with many followers and few friends - and ‘miscreants’ - users with few followers but many friends - are discussed.

Applying the community idea to retweeting, it seems likely that a user in a community is going to retweet posts from other users in the community (due to the shared interests). This is useful when the user has followers outside of the community, since they are then less likely to be following the source user. However, if the retweeter has many followers inside the community, then they are similarly likely to be following the source tweeter and so would be exposed to the original tweet anyway, since there are more interlinking edges between users in the graph. If many users within the community end up retweeting a post, then other users in the community may end up be exposed to the same tweet retweeted several times. This paper does not study this problem, but more analyses the behaviour of retweeting in general, paying particular attention to any trends or patterns arising.

## 4.4 Twitter Analysis

Discuss wanting to know more about Twitter and the way information is propagated through this.

Want to gain a base knowledge of some of the propagation patterns in terms of the tweets and the networks the tweets travel through.

This will form basis of research further on.

Stuff interested in:

- retweet groups (define)

- path-length (define)
- friend/follower graph
- timings, etc.

## 4.5 Experimental Methods

### 4.5.1 Data Collection Methodology

The experiments conducted involve the examination of tweets selected randomly from the Twitter public timeline<sup>1</sup>. The system interfaces with the Twitter REST API and periodically retrieves the top 20 tweets from the timeline, filtering out the posts that are retweets. They are distinguishable as they start with the characters ‘RT’ followed by a username. These retweets are then stored and the original tweet text is extracted. Sometimes retweets are made with additional text as comments, which made it harder to retrieve the actual original text. In these cases, additional queries are made to try and source the original text, but, failing that, the retweet is discarded.

The original text is then used to query all available tweets for the original tweet and any other retweets. The original tweet, along with its set of retweets is referred to as the *retweet group*. As a result, each retweet group contains precisely one original tweet and at least one associated retweet, thus the smallest possible retweet group is of size two. The *final retweeter* is the most recent person to retweet the post. Results are generated by extracting data from these retweet groups, which are stored with metadata such as time, source, language, and so on.

It should be noted that there are two ways in which tweets can be retweeted. The traditional (or *old*) manner involved users copying the existing tweet (or retweet) from the timeline before tweeting that with the characters ‘RT’ and the previous retweeters username at the start. This ensures that the appropriate users are cited as being the

---

<sup>1</sup>Viewable from <http://www.twitter.com/> when not logged in.

author of the tweet. The *new* method enables users to retweet a post simply by clicking the ‘Retweet’ button that is available next to the post on the timeline both on the main Twitter website and on various Twitter clients. Despite the fact that the website and clients have alternate ways of displaying these types of retweets, the search interfaces and API still see them as starting with ‘RT’ and the previous poster’s username. This, fortunately, means that both types of retweets can be collected in the same way. Retweeting in the latter method, however, is unavailable if the post that is to be retweeted belongs to a user who has a non-public profile.

Despite this, the new method doesn’t have direct support for retweet chains (explained below) to be formed. No matter who retweeted the post previously (whether they were the original poster, or not), the new retweeting method simply treats it as a direct retweet from the original poster (as long as the previous retweets were also made in this fashion). The original tweet (and other retweets) can still be found using the search method outlined above, however, so this doesn’t significantly affect results. Limitations with both the Twitter API and the search interfaces did mean that not all retweets could have successful queries made against them.

As briefly described above, these experiments use the notion that retweets themselves can be retweeted; analogous to routers forwarding packets in a multi-hop networking system.

### 4.5.2 Data

There are three sets of experiments detailed in the following section and all of the data was collected using the search method outlined above. The data consists of around 4400 retweet groups (representing a total of about 26,000 tweets and retweets) collected between 26th January & 24th May 2011. It is acknowledged that this is a relatively small result set and so we emphasise that these are *preliminary* experiments to provide insight into the decentralised side of Twitter and to act as a start point for future work


in this area.

## 4.6 Empirical Results

### 4.6.1 Length of Retweet Chains

The longest path-length encountered from the dataset of retweet groups was 9, meaning that the tweet has been forwarded nine times from the original user; making the total number of involved users ten. Despite this, in the dataset, the average path-length was around 1.8, with the vast majority of tweets being retweeted between one and two times, as shown by the distribution in Figure 4.1. The *maximum path-length* of a retweet group refers to the longest path-length present in the group.

It should be noted that the very high proportion of single-chain path-lengths could be



**Figure 4.1:** *Log/log distribution of maximum path-lengths from retweet groups.*

a result of the new method of retweeting introduced in the previous section, since, as mentioned, if all of the users in a retweet chain were to retweet in this way then Twitter would still only treat the path-length as one; between the original tweeter and each of

the intermediate retweeters. Despite this, the result fits the rest of the dataset well.

This data can also be used to have a look at the friend-follower graph associated with retweets. In cases where the group's maximum path-length is equal to one (i.e. simply the case where a user has retweeted another user's tweet once), the retweeter follows the original tweeter about 90% of the time. This implies that in 10% of cases, a retweeting user has retweeted a tweet found outside of their home timeline, i.e. on the public timeline, or whilst browsing through other users. It could also be as a result of the retweeter not including the tweet citation ('@username') for reasons such as to save space, which would be particularly likely in cases of longer tweets. This implies that the more followers a user has, the greater chance that user has of having his or her tweets seen and of being retweeted. Since 90% of retweets occur if the retweeting user follows the original tweeter, then this is directly demonstrated. However, this does not take into account whether the tweet was passed down a chain of retweets retweeted with the second (single click) method. In these cases the path-length would be represented with a length of one even if they were actually longer. The introduction of the second method of retweeting is, therefore, not helpful with the analysis of retweets.

In addition to this, in cases where the maximum path-length is greater than one (i.e. at least one user retweeting in a chain between original tweeter and final retweeter), the final retweeter follows the original tweeter in about 40% of cases. From Figure 4.3 we can see that retweet groups with a larger maximum path-length tend to be larger themselves. This means that the tweet has travelled further both around the original tweeter, but also potentially to other communities. Users from other communities will be less likely to follow this original tweeter, explaining this drop in likelihood. Results have also hinted that the probability of users citing the original tweeter decreases as the path-length of the retweet increases, which also goes some way to explain this point.

### 4.6.2 Size of Retweet Groups


This section focuses on properties of retweet group size, and also how this can be related to path-length. The average retweet group size found from the data was of size 6, with the largest and smallest retweet groups being of size 284 and 2 respectively.

The distribution of retweet group sizes can be seen in Figure 4.2 and shows that this data follows a power-law distribution (with a relatively large  $p$ -value of 0.871). The  $Pr(X \geq x)$  represents the complementary cumulative distribution function, where the randomly generated  $X$  is greater than or equal to  $x$ . These values were calculated using the techniques and code provided by [6].


In these retweet groups, and particularly the larger ones, retweets are likely to be down several chains. Combining this information with that of the average path-length of retweets, it becomes more clear how the retweet tree structure is created, with nodes being the retweeting users (the root node being the original tweeter), and edges between nodes being the path of retweets. Different retweet chains would be illustrated by the combinations of journeys down different branches. The longest path-length would be represented by the depth of the longest tree branch, and the total number of retweets would be represented by the number of edges. The trees can be said to represent the retweet groups, though, in practice, many members of the group will be separate from the main tree due to missing citations.

[12] discusses retweet trees in more detail, in particular, the different shapes of the trees representing retweet groups from the Air France incident. Their results, for these trees, also (see section 4.6.1) indicate a large number of groups with maximum path-lengths of one and two.

Figure 4.3 shows how the total number of retweets varies with the longest path-length of the retweet group. The trend mostly correlates with what might be expected; that the maximum path-length of a group increases with the overall size of the group. These illustrations do not show which users are followers of others, but do show how some



**Figure 4.2:** *Maximum likelihood power-law fit for the cumulative distribution of retweet group sizes..*




**Figure 4.3:** *Log/log relationship between the maximum path-length and size of a retweet group..*

tweets are retweeted significantly more than others. Users who have many more followers are said to be more *influential*, say Cha M. et al. [4], who also discuss the idea of ‘indegree’, and that those users who have far more followers than friends are likely to be far more influential. Their chance of having tweets retweeted is therefore increased.

The (immediate) audience size of a retweet group refers to the number of users the tweet has been exposed to either directly or through retweets. *Immediate* was used in this sense since users who have made their profile public can have their tweets viewed by users who aren't followers of the former. This audience size can therefore be calculated simply by summing up all the followers of each of the users in the retweet group.

These audience measures take into account that some users may be exposed to the



**Figure 4.4:** *Relationship between a retweet group's audience size and its longest path-length..*

same tweet more than once. This happens when users involved in a retweet tree have some followers in common, and so is likely to be prevalent in more closely-knit communities. As a result, the audience sizes shown only represent *distinct* users exposed to the tweet.

Collection of the audience size from the data started slightly later than the main dataset and is only available for 2860 of the total 4400 groups. The longest maximum path-length of this subset is 8.

Figure 4.4 shows how the audience size of retweet groups varies with their maximum path-length. The peak at path-length 5 indicates that the groups with a mid-range maximum path-length tend to have a larger audience size, and we believe this is to do with



the amount of audience overlap in different retweet groups.

[12] discusses how retweeting is related to the audience size of a tweet and how the power of the retweet phenomenon can greatly affect the spread of information, even if the original tweeter has only a few followers. The same paper more specifically mentions that the audience size of a retweeted tweet reaches, on average, at least 1,000 users, no matter the number of followers of the original tweeter. This can also be seen in our results; that no matter the maximum path-length of a tweet, the number of users reached is relatively high.

The overhead of a retweet group represents the number of users who are exposed to a tweet more than once and is present in 71% of retweet groups. The *proportionate* overhead is the ratio of overhead to audience size and the mean of this was found to increase with the group's maximum path-length. This is a possible explanation for the peak in the data: that eventually the overhead of non-distinct users has increased to the extent that it reduces the audience size more significantly. The same graph representing the effective audience size (calculated with the addition of non-distinct users) represents, mostly, a continuous positive correlation.

Three of the largest five overheads collected were from retweet groups with a maximum path-length of 1, the largest with an overhead of 6.5 times greater than the distinct audience itself (the overhead was larger than the actual audience size in around 3% of retweet groups). This shows that, to an extent, there can be significantly more overlap in more closely-knit communities; those retweet trees which are wide and shallow. The chance of getting no overhead increases in smaller retweet groups.

### 4.6.3 Retweet Follower Pattern

This experiment focuses on the pattern of followers in the retweeting hierarchy.


The first result shown from the experiment is that the final retweeter follows the previous retweeter in the chain in 67% of cases. It initially seems strange that this should be 20% lower than when following a user in retweet chains of length one. This sug-

gests that users involved in shorter-chain path-length retweets are members of more tightly-knit communities. Retweets with longer path-lengths have, by nature, travelled further and so would be the type of retweet to travel between communities, reducing the chance of the involved users following each other.

The interesting part of this, however, is the number of followers of the previous retweeter in different cases. In the 33% of cases where the final retweeter doesn't follow the previous retweeter, the latter has, on average, around 600 followers. When the final retweeter *does* follow the previous retweeter, however, the previous retweeter's average number of followers is 940. This is quite a substantial difference and certainly highlights the fact that by having more followers you are more likely to have more influence in terms of whether you get retweeted, or not.

This is accentuated further when looking at the original tweeter. The likelihood of a retweeter following the original tweeter in cases in which the path-length is of more than one has already been found to be around 40%, but the average number of followers of the original tweeter increases by a factor of around four (580 to 2000) when also followed by the final retweeter. Results showed that the original tweeter had a consistently higher number of followers when followed by the final retweeter than when not at all path-lengths. This demonstrates that having an increased number of followers is correlated with the chance of a user being retweeted. In this case, having four times the followers increases the correlation dramatically (40% to 90%). The number of followers of a user can therefore be directly related to the ideas of influence discussed in [4] and also of 'advertising' themselves.

It was found, however, that the number of followers of the original tweeter diminishes as the path-length of the tweet increases (Figure 4.5), signifying that tweets travel further when the original tweeter has fewer followers. Because the retweet groups were collected in such a way so that groups containing longer path-length retweets also contained many shorter-chain retweets, retweet groups containing path-lengths of 5 (or more) are also likely to contain many retweets (if not more) with path-lengths of one or two (see the distribution in Figure 4.1). It can therefore be argued that there are




**Figure 4.5:** *Relationship between number of followers (and respective distribution) of the original tweeter as the path-length increases..*

more users involved in shorter-chain retweets than in ones with longer path-lengths. It is then more likely for these users to have more followers than others in the retweet group. Another explanation could be that users are actually aware of their local network and realise that retweeting may cause a lot of audience overlap (particularly in the case of large communities). A user may have seen a post retweeted a few times on their home timeline and thus decide not to also retweet.

One last interesting point to make regarding the notion of retweet chains is looking at the how the pattern of following previous retweeters develops as the path-length increases. It has already been discussed above how the chance of following the previous retweeter in the chain is about 67%, but, in cases where the path-length of a tweet is greater than two (i.e. at least two intermediate retweeters between final retweeter and original tweeter), the chance of the final retweeter following the next retweeter along preceding the previous retweeter is around 45%. This suggests that retweeting is more widespread and not so much just circulated around communities. These preliminary results demonstrate that the chance of the final retweeter following previous retweeters - up to and including the original tweeter - diminishes along the chain or as the tree is ascended (Figure 4.6).

Because of this, it's sensible to assume that the tweets in the dataset are forwarded



**Figure 4.6:** *Proportion of final retweeters following upstream users at varying distances along the chain..*

through less-connected users, and perhaps forwarded from community to community by those users belonging to several groups. Otherwise, if the retweets were circulated more around closely-knit communities, the likelihood of the final retweeter following the previous tweeters would be both greater and more evenly spread - i.e. the chance of following the previous retweeter would be roughly equal to the chance of following the other tweeters in the chain.

In addition, of the 67% of final retweeters who *are* following the previous retweeter, about 19% of them also follow the next previous retweeter (i.e. the retweeter at path-length - 2). In this case, the next previous retweeter has, on average, 3000 followers. In the 81% of these users *not* following the next previous retweeter, then the latter has an average of 525 followers. This is an accentuated result of the one previously, but this time boasts an increase of a factor of 6.

Of the 33% of users who *don't* follow the previous retweeter, about 30% follow the next previous retweeter. Both of these sets of statistics also go towards the idea of the diminishing chance of following the users as the tree is ascended.

From this dataset, it was also possible to work out how often retweeters cited the original tweeter of a post. In retweets, users are typically cited by, as we have seen, having their name along with an 'RT' at the start of the post. This data was collected by seeing


if the original poster's username was mentioned *anywhere* in each retweet. The chance of this occurring was found to be around 68% and did not vary with any pattern with path-length.

#### 4.6.4 Retweet Time Delay

The final experiments in this section focus on the time delay between the final retweeter and original tweeter. This is an interesting area since it enables researchers to see how fast messages propagate through the Twittersphere. From this information, and by using the retweeter patterns demonstrated above, it would be possible to work out how far and how quickly information can be passed around.

Figure 4.7 shows the average time delay between the first and final retweet with increasing maximum path-length of the retweet group.

The results indicate that, mostly, as the group's maximum path-length increases, then



**Figure 4.7:** *Average time (in seconds) between first post and final retweet of a retweet group varying with the group's maximum path-length..*

so does the elapsed time between the original post and the final retweet. This is probably as was expected, since this shows that it takes longer for a retweet to travel further. The data is not consistent, however, especially results for a path-length of five and above. The first four results suggest a uniform incline roughly proportional to  $v = \frac{s}{t}$ ,

where the distance,  $s$ , is the hypothetical distance given by the path-lengths, showing that the *speed* of propagation remains mostly constant.

There is not enough of a trend in the data to make any deductions regarding propagation speed, however. There are two main conflicting arguments regarding this result: the first is, as mentioned, that the further a tweet travels the longer time it travels for. The second is to do with tweet popularity: the more popular a tweet is, the more quickly it will be retweeted. In the latter case, it is possible that longer trees grow fully before shorter ones, implying exponential growth. Generally, though, it seems that the maximum path-length of a retweet group does not massively affect the tweet's propagation speed.

## 4.7 Summary

The experimental results have certainly highlighted the ideas of communities and that of message cascading similar to that demonstrated in [11] and [8] respectively.

As has been seen, in section 4.6.4, the retweet tree seems to grow in a variety of ways. One argument is mostly expected; that as the 'distance' the tweet travels increases, then so does the time taken for it to reach its end. The other argument is linked to the idea of tweet advertising, discussed in previous sections, and namely the notion of tweet popularity.

The previous experiments showed how the number of followers of a user directly influences their chance of being retweeted. It can therefore be seen that advertising can also be linked to the level of influence a user has. The results help illustrate the multi-dimensional properties of the retweet tree, and how these factors relate to its growth and its associated friend-follower graph.

This paper has demonstrated relatively simple results in an attempt to realise some of the behavioural patterns of retweets, both linking to the psychologically in terms of the users, but also the physical properties of retweets. The results are able to repres-

---

ent a basis for potential further work and research into the various aspects of Twitter, moreover, perhaps, topical categorisation and the dynamicity of the friend-follower graphs.

---

## Chapter 5

# Analysis of Twitter's Social Structure

Stuff to add to this section:

- Change tweet features for each simulation and make comparison on these differences
- Observe differences in patterns when network generation parameters are altered
- Link up section to previous section (i.e. how did the previous research help and how does this build on that work?)
- Explain how this section becomes the basis for work in 'main chapter 3'.
- (e.g. Issues with current method (too long, requires network, inaccurate due to having to choose users with fewer followers), so need a quicker, more accessible and online approach).
- Explain the Mechanical Turk questions in more detail, with examples.
- Discuss about the machine learning approach used (logistic regression and how it works)
- Link 'retweet volume' to 'retweet group size'

Twitter is often seen as one of the biggest sources of new and live information on the Internet, with millions of people producing and absorbing information daily. Users



receive tweets onto their timelines from the users that they follow. Thus, a user has some control over the *type* of information they receive by choosing which other users to follow. A particular user, therefore, may not be aware of information that exists outside of their local network, since he or she is not directly exposed to the information produced by non-followees.

Retweeting allows users to forward information they receive onto their own followers. As a result, followers of these users, who might not normally be exposed to this information, now have a chance to access it. A user who decides to retweet a tweet can be said to consider that tweet to be *interesting* (at least, to their followers), since that user has taken the time to read the tweet, decided whether or not to share it, and then to actually retweet it [16].

Several factors can affect a user's decision to retweet, such as whether the tweet contains an (interesting) URL, whether the tweet mentions another user, whether a user even has a chance to see the tweet, the influence of the author, and so on. These factors account for a user's individual retweet *decision* on a particular tweet, and the combination of several users' retweet decisions dictate how far the tweet will propagate. However, it is our belief that the social network structure also has an affect on how far tweets can travel.

The social structure of Twitter is built up by users electing to follow other users. When a user follows a user, a directional link is forged between them, and any tweets generated, or forwarded, by the followed user are passed down the link. It's clear to see that, as more links are made between users, many more avenues are generated for message propagation throughout the social structure. Users with a high in- and out-degree can become an information highway, but users with a low out-degree are a bottleneck of information.

In this chapter, we demonstrate how different network types support different propagation characteristics through the use of a model simulating each network type. Using the model, we make predictions on the retweet outcomes on several network types, and compare these to the characteristics of the propagation in real Twitter networks. We

finish by discussing how the *interestingness* of a tweet may be inferred from simulating the network in this way.

## 5.1 Background to This Area

[16] discusses an idea similar to that of predicting the interestingness of a tweet, but focuses mainly on predicting the users most likely to find the tweet interesting enough to retweet. Similarly, [10] looks at the same issue, but at the other way around by predicting the *type* of tweets that are likely to be retweeted many times. Discussions on the retweet decision in relation to a user's recognition of its features take place in [5] as well as conclusions about the effect of features such as the number of followers of a user or any pre-existing metadata on the interestingness of the tweet.

The notion of time decay and how this is associated with a user's retweet decision is discussed in addition to a retweet probability prediction in [20] (and [14]), which is the basis of the model we use in this paper.

## 5.2 Overview

In the next sections, we introduce and briefly explain the regression model we use for simulations and predictions. We then go on to analyse the differences in the propagation characteristics between three different network types before comparing the results of simulations on these networks with data collected from the Twitter social graph. We finish by introducing a methodology for predicting the interestingness of a particular tweet to a particular user and how this might be improved.

Ideally, there are two things we'd like to see from the first set of experiments; firstly, that changing the network type and properties does, indeed, affect the propagation behaviour, and, secondly, that at least some of the results from the experimentation correspond to Twitter's own retweet behaviour so that a fair comparison and justifica-

tions of our results can be made.

With regard to our prediction work, we'd like to be able to make relatively decent predictions on which tweets are of interest, and which are not, based on the simulation research in the next section.

## 5.3 Model

As mentioned above, [20] introduced and discussed their prediction model, which was shown to perform well when predicting the retweet decision of a user. Their model trains a logistic regression using a set of user-, tweet- and context- features in order to classify an experimental tweet and output a retweet probability based on a set of features.

### 5.3.1 Machine Learning

Machine learning techniques are useful for making predictions based on certain input criteria based on the perceived history of previous outputs from the same inputs.

#### General Overview

For example, consider the three attributes, A, B and C, each of which is of a boolean data type. A machine learning technique is 'shown' that in every case where A is True and B is False, then C is true; and that in each case where A is False and B is True, then C becomes False. The history of these inputs suggests that A is strongly associated with C (the strength of this relationship will increase if the input/output history is larger), such that if the technique now tries to predict C based on the fact that A = True, then it will likely suggest that C is also True with high confidence.

A technique that has been shown many of these input/output combinations (known

as ‘instances’) is said to be a trained model, where the model type is that of the learning technique used. This model can then make predictions based on the same inputs it has been trained on, and will therefore not work for attribute inputs that it hasn’t been trained with.

Generally, in large enough datasets, the training data will not exclusively contain instances where  $A$  is the inverse of  $B$ , and so the model will be able to make predictions in cases where  $A = B$  (though if these cases are less prevalent then the confidence of these prediction outputs will be weaker). When training the model, instances must contain the input attributes as well as the value of the output attribute. When testing against the model after training, the model is supplied the input attributes and predicts the value of the output attribute. Generally, the model will also indicate how confident it feels that the prediction is correct, and thus an idea can be obtained of the value of each of the input attributes. Some model types (particularly regressions) output instead the *probability* with which the attributes align to the trained regression.

In this example the focus has been on boolean data types, but most machine learning techniques are not limited to these. Indeed, most applications require the learning of real numbers (integers, floats, etc.) and more highly-dimensioned nominal attributes, where there are several categories that the value of one of the input attributes can lie in. Many machine learning techniques exist and some are more accurate than others when it comes to different data types and several support the notion of ‘weighting’ attributes, in which a strength weight is assigned to each attribute to signify the confidence with which the model should rely on that attribute for producing the prediction.

For example, for purely nominal values, then logistic regressions can be accurate in predicting outcomes, whereas for a mixture of real and nominal attributes a Bayesian attempt might be more suitable.

## Logistic Regression

The work in this chapter utilises a logistic regression for outputting retweet predictions based on the feature input criteria discussed later.

As mentioned, the logistic regression predicts the statistical likelihood that the input features align with what the model has been trained with. The work in this chapter relies on the use of a retweet probability of a given user acting on a given Tweet at a certain time. Thus, it is trained in such a way so that the binary input attributes produce a binary output retweet outcome attribute: 1 for retweet, 0 for no retweet.

When testing, the binary features of the test set are used to output the probability that these features will result in a retweet.

### 5.3.2 Algorithm

Include image indicating the flow of events that take place for the model

Pseudocode for the algorithm

In essence, the model requires a network of users and one tweet to start the simulation. It starts by initialising a user set,  $U$ , to contain one source user from the network,  $U_s$ . This source user then transmits a tweet which is then received by its followers.  $U_s$  is then removed from  $U$ . For each follower, its tweet and user features are classified to produce a retweet probability. If this is greater than a random number,  $R$ , then the user retweets the tweet to its followers and the process repeats. Any user that retweets the tweet is removed from  $U$  and added to the retweet decision set  $RT$ . The followers of the retweeter are added to  $U$  for the next iteration, in which each user in  $U$  has received the tweet onto their timeline. The retweet volume of this tweet in this network is then the cardinality of the set  $RT$ .

### Twitter Timelines

A tweet in a user's timeline will slip further down in the timeline as time goes by. This happens whether the tweet is interesting or not and whether or not the user has even seen the tweet. Users having a quick browse through Twitter may not have time to scroll down to find these interesting tweets (and will not know they exist) and thus tweets left in the timeline have their retweet chance decay over time.

We emulate this phenomenon in our simulations by removing the tweet from a user's timeline (by removing the user from the set  $U$ ) if the user has not retweeted the tweet within a timestep threshold, which can be varied to alter the volatility of the retweet.

#### 5.3.3 Features

We use this model as part of a simulator in order to obtain an approximate retweet volume of each tweet when the retweet probabilities of each user receiving the tweet are combined. [20] used a set of around 50 features, though they mention how some features have more weight than others. In order to simplify the simulation significantly and to make data collection more tractable, we decided to use only the following four major features;

- *follows* - Whether or not the user exposed to the tweet is a follower of the author;
- *followed* - Whether or not the user exposed to the tweet is followed by the author;
- *mentioned* - Whether or not the user exposed to the tweet is mentioned in the tweet;
- *URL* - Whether or not the tweet contains a URL;

Where *author* is defined as the user who originally tweeted the tweet.

## 5.4 Training the Model

### 5.4.1 Data Collection

To collect the training data, we crawled Twitter, using the REST API, from March-June 2012 to collect a set of around 12,000 tweets and retweets from the Twitter public timeline. In addition, we made further calls to collect the information required for the features (namely the *following* and *followed* features).

### 5.4.2 Feature Extraction and Regression Training

FILL THIS IN

The above features were extracted in each case and the regression was trained.

## 5.5 Network Analyses


In this section, we look at three different network structures and discuss the differences in the propagation patterns produced by each. Note that each graph we assess is *directed*. The same set of tweets are used for each simulation on each of the networks.

### 5.5.1 Path Network

A path network is the most simple of the three, and is also the least life-like when compared to Twitter's own social graph. In this network, the output retweet volume is, by definition, equal to the penetration (i.e. *depth* of propagation) of each tweet.

In this case, a directional path network consists of a network of users,  $N$ , of size  $n$ , in which each user  $N_i$  is followed by user  $N_{i+1}$  for  $1 \leq i < n - 1$ . As a result, all users in the network, except the user  $N_n$ , have precisely one follower.

Since each internal user only has one follower, the likelihood of a retweet occurring at each timestep is somewhat reduced, it is expected that the retweet volume will tail off more soon than in other network types. Propagation is also hindered by the fact that each retweet can reach an audience with a maximum size of 1 at each stage, thus relying on that single user making the retweet. Figure 5.1 shows the frequency distribution for



**Figure 5.1: Retweet volume frequency distribution from path network simulation.**

a path network when simulated with the logistic regression over a series of tweets. The graph shows a very large proportion of single retweets, which reduces logarithmically with larger volumes.

The likelihood of a user getting the chance to retweet, and also deciding to retweet, becomes the product of a probability function the further the tweet travels, where user  $N_i$  requires all users  $N_0$  to  $N_{i-1}$  to pass on the message before it even gets a chance to make the retweet decision.


As a result, if the retweet decision chance of each user is more or less equal, the chance of user  $N_2$  retweeting the tweet is of an order of magnitude less than that of user  $N_1$  deciding to retweet. The graph shows half life-style behaviour; owing to the fact that each retweet is exponentially less likely to occur than the previous retweet. Note that the graph is plotted on a log-linear scale.



### 5.5.2 Random Network

Random networks are more similar to Twitter's own social structure than path networks, but are a much more basic and uniform version and do not consider more influential users or the development of Twitter communities.

A random network is defined as a network of users,  $N$ , of size  $n$  in which a user  $N_x$  has probability  $p$  of following user  $N_y$ . Thus; as the probability  $p$  is increased, the likelihood of a user following other users in  $N$  increases, causing the overall network edge density to increase. Generally, the average number of followers and followees of a user is proportional to  $p \times n$ . Thus the parameters for constructing such a graph are the network size,  $n$ , and the attachment probability,  $p$ . The simulation results for the random network indicates a higher distribution of mid-range retweet volumes.



**Figure 5.2: Retweet volume frequency distribution from random network simulation.**

### 5.5.3 Scale-Free Network

Include more mathematical analysis of scale-free networks (in general) - i.e., in what way are they logarithmic?

A scale-free network is a network of users,  $N$ , of size  $n$  and is generated in such a way so that the resultant distribution of degree follows a power-law. *In*-degree signifies the number of inward edges to a node (i.e. the number of followers of a user), whereas *out*-degree is the number of outward edges (i.e. the number of users that user follows). Scale-free networks have been the subject of a fair amount of research, and are explained more thoroughly in [9]. In our implementation we use NetworkX<sup>1</sup>, a Python networking package, to generate directed scale-free networks through a preferential-attachment algorithm based on the network size and edge density as parameters.

Figure 5.3 shows the frequency distribution of retweet volumes. Since the data is plotted on logarithmic scales, we see a logarithmic trend very similar to our results in [17].

#### 5.5.4 Comparison to Real Twitter Data

In our previous work, [17], we captured and analysed data which contained results on the distribution of retweet group sizes. In that paper, a retweet group was defined to be a set of tweets containing one tweet and then all the retweets of that tweet. Thus the retweet volume looked at in this section is effectively the cardinality of the retweet group (minus one). Since the results in the above experiments also look at the frequency distribution of retweet volumes, then we should be able to draw some comparisons.


We compared the data produced by the different types of network to this previous data, and found that the scale-free network produced a distribution similar to that from the real Twitter data.

#### 5.5.5 Structure Comparison

Each network structure has been demonstrated to show different propagation characteristics. This has shown that, in addition to a user's own retweet decision, the actual

---

<sup>1</sup><http://networkx.lanl.gov>



**Figure 5.3: Comparing the retweet volumes distribution from scale-free graph simulation to data from Twitter’s graph.**

spread of a tweet depends somewhat on how the author’s local network is constructed. With lots of edges in the graph, there are many more paths down which propagation can occur, increasing the number of times a retweet decision is made, and therefore an increase in the overall number of retweets occurring. The retweet decision facilitated by the model, therefore, combined with a user network give an overall *retweetability* of a tweet that will vary depending on the network it’s being propagated in, the source user, and the intermediary retweeters.

The path network, as designed to be the extreme case, has shown to allow poor propagation. Importantly, whilst the network parameters and retweet decision had to be globally increased to obtain any sensible data from this simulation, the trend still shows how propagation down a single chain isn’t hugely effective.

The random network facilitated many more retweets due to the fact that users had a very similar in- and out-degree in all cases across the network. This means that each user is able to receive a lot of information, and is also able to pass on (whether an author or a retweeter) information to lots of users simultaneously in the graph.

Despite random networks supporting large retweet throughput (i.e. high *recall* of information), the disadvantage is that the interest *precision* is much lower. This is be-

cause this type of network relies on users following a large number of other users, thus meaning that they would receive more ‘noise’ (i.e. uninteresting tweets) than if they were more limited and selective. Although tweets that are retweeted are usually of a higher *quality*, not all retweeted tweets will be interesting to all users.

Finally, the scale-free network, whilst not having the highest throughput of information, does have trends most similar to the data on retweet distributions collected from Twitter’s social graph. This is due to its ability to emulate more influential users and areas of dense communities (as discussed in [11]). These networks have the potential to allow for large numbers of retweets, especially if they are sourced from one of these more dense areas, but typically die off more quickly as the tweets are retweeted through less influential users.

## 5.6 Interesting Additional Findings

### 5.6.1 Graph Density

Define edge density (with equation

### 5.6.2 Results

Show links between follower audience -> local network -> density.

### 5.6.3 Uses

From this data, it is demonstrable that various parameters can be generally successfully inferred from very basic user information.

## 5.7 Predictions From the User Graph

The final part of this paper focuses on the ongoing development of a method to predict the interestingness of a tweet based on the work in the previous sections. The prediction method, at a high level, compares the predicted retweet outcome of a given tweet to the number of times that tweet has actually been retweeted. If, for example, a tweet is simulated with the help of the model and produces a prediction of two retweets, but the tweet has actually been retweeted four times, then we can infer that this tweet is more interesting (at least, to a subset of users).

The tweet and user features we looked at earlier in this paper are very static, binary features, which do not take into account the actual content of the text of the tweet. Therefore, if a tweet is retweeted more than was predicted, then there is something in the tweet, such as a link to a particularly interesting article or a breaking news story, that makes it more interesting than the average tweet, with the same binary features, that was used to train the model.

In order to improve the fairness of the experiment, we wanted to ensure that the environment of the tweets (i.e. the user network they are propagated through) is the same as its real-life Twitter counterpart. We could then choose a user, which would become the source user,  $U_s$ , in the set  $U$ , and simulate that user's own tweets within their particular local network as described by the model above. This would then produce a retweet volume for this user's tweets, to which we could compare the number of times that tweet has *really* been retweeted.

### 5.7.1 Data Collection

Due to the exponential scaling properties of Twitter's social graph, it was infeasible to collect any more than two hops away from each user as a representation of that user's local network under the rate limitations of Twitter's REST API.

In particular, a single Twitter account running an instance of an application was al-

lowed, at the time of these experiments, a maximum of 350 REST API calls per hour. One call would be required, for example, to obtain up to 5000 of the followers of a particular user (i.e. one follower hop from the user). An additional call would then be required to collect each of that user's follower's followers (in order to obtain the *second* hop from the source user).

Thus, a user who has 700 followers would require 700 API calls to collect that follower network, in addition to the one required to collect that source user in the first place, and would therefore take over two hours of collection. To collect the *third* hop from the source user would drastically multiply the number of required requests (even if there is significant overlap between the followers) and the time needed.

If each of the 700 followers of the source user had, on average, 200 followers, then this would require the gathering of  $700 \times 200 = 140,000$  users, equating to more than 402 hours of data collection. Bearing in mind that this would only collect the network features for *one* user, it is clear to see how this is an impractical approach.

Luckily, in [17], we found that the vast majority of retweets occur *within* two hops of the source user (i.e. a path length of less than three), so we considered that the distance from the source user in each case would be sufficient.

In June 2012, the Twitter REST API was used to conduct a random walk through the social graph. For each user, we collected the most recent 300 tweets (including each tweet's metadata - particularly their retweet count) and their local follower network within two hops. We didn't collect the friend network, as we were only interested in tweets propagating outwards from the source user.

After processing that user, the walker chose a user at random from the present user's set of followers and made this the new current user from which to collect data for. If the present user, at any stage, does not have any followers, a list of previously accepted users is maintained and a follower is chosen from one of those instead.

The walker continued until the rate limit was met, at which time the current state was written to disk, and the walker waited until the rate limit was reset before continuing.

Generally, this resulted in, for each user, a set of up to 300 tweets (totalling to around 10,000 tweets in total) and the network in which these tweets were propagated within. There was no need to collect any further data to train the regression, since we were able to re-use the trained model we used earlier.

### 5.7.2 Validating Results

Needed to validate results using human input. Machines themselves are generally unable to express human interests, so results need to be properly evaluated.

#### Crowdsourcing

Discuss crowdsourcing, its uses, how it is useful in this area. Talk about its history (with any references), and then about mechanical turk.

Mention mechanics of mechanical turk, how it is US only (but we used crowdflower - which automatically handles submission to MT and several other crowd-sourcing services).

#### What We Wanted to Assess

Used MK, etc.

#### Constructing the Questions

Set up questions (i.e. 5 tweets - choose most interesting and least interesting), give example of this.

In order to validate our prediction results, we ran a pilot user study in order to obtain some human input on the interestingness of each tweet. We compiled the tweet data into a set of questions which were submitted to Amazon's Mechanical Turk. Each

question consisted of five tweets from our dataset and each Mechanical Turk Worker (MTW) undertook five questions. Each question asked the MTWs to select which tweet was the most interesting of the five, and which was the least interesting.

For consistency we ensured that at least three MTWs had answered each question. When selecting tweets to include in the Mechanical Turk questions, we excluded those which are ‘@-replies’ - i.e. tweets which begin with another user’s screen-name and typically form part of a conversation between two or more users. This meant that there were around 4,500 tweets in total in the questions.

Through using the model and simulating each user’s tweets through their individual local networks we achieved around 86% accuracy in correctly predicting the number of times each tweet was retweeted.

The precision in predicting the *interestingness* of each tweet was around 30%. While this value is low, it does mean that in 30% of cases, a tweet that we predicted to be interesting was verified to be interesting by at least three MTWs all selecting one tweet from a set of five. In addition, when simulating the questions by randomly choosing the most ‘interesting’ tweet of the five in each case, the performance was unable to near our precision even after several thousand iterations.

### 5.7.3 Improving This

Need offline methodology.

One route for this would be to try and infer a user’s local network from a set of their immediate parameters, drawing on our earlier work suggesting that the Twitter network has the properties of a scale-free small-world graph. Through studying graph patterns, it is possible to make sensible inferences on the edges and nodes of a user’s local network based on their follower count. From this, a graph edge density can be calculated,  $d = \frac{|E|}{|N|(|N|-1)}$ , for use in generating a scale-free network.

Since, for these preliminary experiments, we were only able to collect data from users with a more modest local network, the real and predicted retweet values were both



relatively low, allowing more room for error. When simulating much larger local networks involving many more real retweets for each tweet, predicting interestingness, with some threshold value, may become more accurate and thus help improve the precision. The reason for this is that the retweet count of tweets that naturally get retweeted many tens, hundreds, or more times is likely to vary more with interestingness than those that are naturally only retweeted very few times.

## 5.8 Future Work

There is much further research that could be carried out based on the results in this chapter. Now that the foundation has been laid for simple retweet prediction based on network analysis, research could begin to look at ways in which, as mentioned, networks could be generated based on a few environmental features surrounding users. This would allow for quick generation of user networks (bypassing the need for data collection) and would also support the same calculations for more highly influential users (users with more followers and more retweets per Tweet).

For this research, the notion of the network will continue and form the basis of the environmental features in the next chapter. Since we now know that the network plays an important role in dictating the way in which information can propagate

## 5.9 Summary

In this chapter we aimed to carry out a study on the behaviour of propagation through different types of social graph structures and to introduce our ongoing work into predicting the interestingness of tweets from their retweet patterns.

Using a set of tweet and user features, we trained a regression model which we used to simulate a number of tweets through different network types. We produced a dis-

tribution of retweet volumes for each network type and confirmed that, with the same tweet features, different network configurations do indeed facilitate different retweet behaviours in terms of propagation spread. We were also able to compare our results to data from Twitter to verify that Twitter's own social graph most closely resembles a scale-free small world graph.

We then finished by discussing how we used the trained model to simulate real networks from Twitter, along with the tweets that were passed through these networks, in order to try to predict how interesting a tweet is based on its retweet patterns. While we were able to often correctly predict the retweet outcome of a tweet, we found that more work would be required to improve the performance of predicting whether or not these tweets are truly interesting to users.

---

## Chapter 6

# Inferring Interestingness of Tweets based on Information Flow Through the Network

Mention:

- How this chapter builds upon network stuff in previous chapter
- We hope to compare and contrast two better ways of predicting retweet volume *and* interestness
- What needs to be improved (speed, usability - more users with more followers etc.)
- Why do improvements need to be made?
- How is this useful, and how does first chapter relate to work done here?

Discuss:

- Does not use network to simulate tweets - instead uses a set of user features
- Previous chapter shown how basic features can be used to generate a scale-free network, which is what twitter is

- Use these features as input attributes of a new machine learning technique model.
- This method does not use a network or model individual user decisions
- Trained on a set of that particular user's tweets with the retweet outcome of integer type
- A new tweet modelled with the regression outputs a retweet volume prediction without having to simulate the Tweet's travels through the network.
- Discuss about the machine learning approach used (logistic regression and how it works)
- Talk about the 'binning' of retweet outcome volumes and its approaches (distribution dependent / independent, tables of precisions, etc.)
- Link 'retweet volume' to 'retweet group size'

Mis-calculation errors when validating the interestingness predictions: - caused by non-authorisation to collect the data from Twitter (i.e. user has a protected account). Therefore we cannot build the test (or train data) successfully for these tweets - if an experimenter selected one of these Tweets, then we have to discard that timeline from our analyses.

In this chapter, we use the features surrounding both the social structure of Twitter and the Tweets that propagate within it to develop a methodology for deciding upon *which* Tweets may be of specific interest, and also an introduction to inferring *how* interesting the information is. The work does not take into account the relevance of a piece of information to a certain user, but instead the general interestingness level of the Tweet. The work focuses on the difference between the raw popularity of a Tweet, demonstrated by its retweet activity, and how interesting the Tweet actually is to its recipients. While it has been shown that training a model with many Tweet features can be accurate in predicting how many times a Tweet might be retweeted [20], here we focus on

the notion of Tweet content beyond those static features. That is to say, when comparing the popularity of Tweets, that there is some content (either in the Tweet itself or in a webpage or image at a URL included in the Tweet) that makes that particular Tweet stand out and to cause *affective stimulation* [18] to viewers.

The focus of this work is in quantifying the interestingness of Tweets; that is, the universal relevance level of the information contained (or portrayed through links or media) in a Tweet. We infer this interest level by analysing the *retweet volume* of Tweets in addition to their surrounding static features and the features of its source user and the network it propagates within. The retweet volume of a Tweet is defined as the total number of times that particular Tweet has been retweeted.

This paper considers only Tweets that have been retweeted using the ‘button’ method, which is a single-click function on the Twitter website and its applications to carry out a retweet. Twitter users also sometimes imitate a retweet by copying the text of the original Tweet and prepending it with an “RT” followed by the original author’s screen-name, which allows them to add annotations to the Tweet if they desire.

In this work, we introduce a method of calculating if a Tweet is interesting, or not. This is done on the basis of a comparison between a predicted retweet outcome for a particular Tweet and the *observed* retweet volume of the Tweet. The method relies on the fact that the content of a Tweet is an important factor in retweet decisions, but avoids having to do much processing of the Tweet’s text contents themselves (such as through natural language processing or sentiment analysis).

To determine *how* interesting a particular Tweet is, we introduce an ‘interestingness score’, which is based on the quantitative *distance* between the real and predicted retweet volumes for a specific Tweet. The method and scoring is discussed in more detail later. If the interestingness score for Tweets can be determined and is known for a set of Tweets, then this method could be used as the basis for an information retrieval or delivery system, where relevant and interesting information can be shown to users without them having to know about it (e.g. follow the right set users) or search for it first.

### 6.0.1 User Influence

Of importance to this work is the difference between the retweet volume and the interestingness of a Tweet, and the fact that one does not necessarily indicate the other. The most obvious case in which this occurs is where there are influence discrepancies between users. For example, one of the most influential Twitter celebrities is Justin Bieber (@justinbieber), who, at the time of writing, has nearly 40 million worldwide followers and achieves an average of around 50-120 thousand retweets per Tweet. His Tweets currently rarely receive less than 40,000 retweets.

Average Twitter users generally attract a couple of hundred followers and would normally receive very few (if any) retweets per Tweet. A particularly interesting Tweet from such a user may be retweeted around 5-20 times (though this depends entirely on the level of influence of a user and their number of followers).

It is therefore apparent that an uninteresting Tweet from Justin Bieber could get retweeted 50,000 times and an exceptionally interesting Tweet from a less-influential user may get 30 retweets, and that the retweet volume, in this case, really isn't indicative of the *interestingness* of the particular Tweet.

Of course, this brings about the notion of information relevance, and the fact that a Tweet could be very boring or irrelevant to one user, and very interesting to another. In this work we focus on *global* (or 'average') interest, where retweet scores and predictions are made for the general case. It is our belief that Tweets that are retweeted more than expected within their authors' local networks (relative to the authors' own regular output) are also likely to be of interest to a wider audience.

## 6.1 Interestingness Scores

As discussed, our proposed process for deciding upon a Tweet's interest level involves conducting comparisons between the number of times a Tweet has been retweeted and the *predicted* retweet volume of the Tweet. The premise behind this is that there

is something beyond the static features of a Tweet (e.g. whether, or not, the Tweet contains a URL, the length of its content, the influence of its author, etc.) that causes a specific Tweet to be more interesting than another with precisely the same feature set. The retweet count of a specific Tweet is returned, at the time of writing, as part of a standard call to v1.1 of the Twitter REST API, and includes the total number of times that the original Tweet was retweeted using the button retweet method.

Essentially, our method follows these basic steps:

1. Obtain and extract contextual information on a given Tweet,  $t$ ,
2. Generate predictions on the retweet volume for  $t$ ,
3. Compare the predicted values to the *observed* retweet count of  $t$ ,
4. Make inferences based on this difference between the predicted and observed values.

We use Bayesian Network classifiers, trained on a set of Tweet and environmental features, in order to make predictions on a Tweet's retweet volume. In this work we consider two of these classifiers for making predictions for each Tweet; one trained on the Tweet's source user's own Tweets and another trained on a 'global' corpus of Tweets (a larger dataset of Tweets from many source users). Thus, we obtain two retweet volume predictions for each Tweet.

For a Tweet,  $t$ , we define the observed (or 'real') retweet outcome to be  $T_O(t)$ , the predicted outcome using the global corpus model to be  $T_{P_G}(t)$ , and the predicted outcome using  $t$ 's author's user-centric corpus model to be  $T_{P_U}(t)$ . The scores for  $t$ 's actual outcome compared to the predicted outcome using the global and user-centric models,  $TS_G(t)$  and  $TS_U(t)$  respectively is calculated thus;

$$TS_G(t) = \frac{T_{P_G}(t)}{T_O(t)} \quad TS_U(t) = \frac{T_{P_U}(t)}{T_O(t)}$$

This provides a positive score where;

$$TS_G(t), TS_U(t) \begin{cases} > 1 & \text{indicates } t \text{ is interesting} \\ \leq 1 & \text{indicates } t \text{ is non-interesting} \end{cases}$$

$|TS_G(t) - 1|$  represents *how* interesting (or non-interesting)  $t$  is and  $TS_{Avg}(t)$  denotes the mean of its user and global score.

This method relies on collecting retweet data from Twitter and involves taking a snapshot of Tweets at one stage during their lifetime. Since Tweets do not decay (unless they are deleted by their author), they can be found and retweeted by users at any time after their composition. In this work, we assume that the significant portion of retweet activity has already occurred for the Tweets that have been collected. Indeed, [12] carried out temporal analyses on retweet behaviour and discovered that, on average, 75% of the retweets of a particular Tweet occur within one day and that 50% of retweets take place within one hour of the Tweet being posted.

Based on this, to ensure that the retweet count mostly reflected the ‘final’ retweet count, only Tweets more than one day old were considered for experimentation.

### 6.1.1 ‘Twitter is a Memepool’

The term ‘meme’ was defined to be a “unit of cultural transmission” [7] and the notion of memetics is an analogy to biological genetics. Unlike genes, memes are entirely non-physical and represent a cultural idea or other human-based behaviour.

Genes survive and are passed on through *replication*, and this replication occurs more frequently and efficiently when they are more suited to their environment. A gene’s genome is the set of information that represents, in its entirety, the features that make up that particular gene (such as eye colour, height, some aspects of personality, and so on).

Memes are similar in that they contain a set of features, such as the wordings of a phrase or their relevance to other cultural aspects, which cause them to be less or more likely to be replicated in different environments. For example, an Internet meme relating



to the Star Wars movies would likely have a greater chance of reproduction (through reposting, discussion, etc.) in an environment of sci-fi fanatics than when amongst more mixed groups.

The meme is a useful analogy for the work in this paper, since it also helps outline the way in which Tweets are replicated within Twitter. Like a meme, a Tweet has a specific set of features (the text it contains, any hashtags, the inclusion of mentions, and so on) and it exists within an environment consisting of a set of interconnected users on the Twitter social graph. A particular Tweet would generally have a greater chance of being reproduced, through retweeting, in certain networks than others. Tweet features are similar to the *genome* of a gene and the features of the network the Tweet exists in (i.e. the users that receive and have the opportunity to assist in a Tweet's propagation) form the *environment*.

More:

- Gene is a physical entity containing information and instructions. It is a unit of genetic inheritance (i.e. offspring typically have a mashup of the genes of the parens)
- The result of the data held by a gene (the genome) means that organisms with certain genes are able to reproduce and survive more than other organisms containing different genes.
- Thus the gene is able to replicate under certain gene- and environmental-centric conditions.
- A meme is similar to gene but is non-physical. They are a unit of cultural inheritance (an idea, phrase, behaviour, etc.).
- Like genes, memes are able to survive better when their features (*menome*) are suited to the meme's environment. In such environments, the meme is able to be shared and replicated more efficiently and frequently.

- A Tweet, again, is similar to both. A Tweet itself has many features (the text of the Tweet, the time of its origin, its length, etc.) and their environment, the Twitter social structure, has features (namely the users that belong to it and the way they are connected) which may facilitate the replication (i.e. Retweet) of the Tweet.
- A Tweet existing in different social structures will have different Retweet patterns, which is what we want to show in this chapter.
- Thus  $\text{tweetfeatures} = \text{genome}$ ,  $\text{userfeatures} = \text{environment}$

## 6.2 Retweet Volumes as Nominal Attributes

In order to try and improve the accuracy of the model at predicting retweet volume outputs, a nominal output attribute would be better than a real one. Predicting a continuous numeric value could render inaccurate results and calculating the cut-off points at which to mark as the upper- and lower-bounds for the output based on the inputs would raise difficulties.

Instead, the retweet outcomes were to be ‘binned’ into several categories which would be determined on the fly based on the outcomes present in the training data.

In the global model and each of the user models, the retweet volume was trained as the *outcome* feature. Accurately predicting continuous values is a challenge for many machine learning algorithms due to the way the model is built around the features. To assist with this in our experimentation, the retweet outcomes were ‘binned’ into nominal outcome categories.

The binning algorithm needed to be dynamically based on the distribution of retweet outcomes so that each bin was of a roughly equal size.

### 6.2.1 Binning the Retweet Outcomes

Describe the different methods (linear, distributed 1, distributed 2), and their advantages/disadvantages, with examples showing the graph of what the bins look like.

Focus on the distributed 2 example, and why this is better. ‘Requested’ bin number not usually the same number as what is actually returned (due to large numbers of smaller retweet groups).

Pseudocode

In [17] it was shown how the general distribution of retweet volumes forms a long tail with a very high proportion of single retweets and for which the frequency drops off considerably for larger retweet volumes. As a result, linearly binning the outcomes would reduce the number of feature instances for higher retweet outcomes and thus cause the Bayesian Network classifiers to be considerably less accurate.

The binning algorithm that was used involved calculating the projected size of each bin from on the total number of feature instances and the desired target number of bins. The bins were then filled accordingly, such that each retweet volume frequency would only appear in one bin. For example, The feature instances with zero retweets would all appear in the first bin, no matter how many there were. Instances with larger retweet volumes would then typically share a bin with instances with similar retweet volumes. Each bin ‘range’ (i.e. the range of retweet volumes of instances contained in the bin) represents a nominal category of the retweet volume outcome in the Bayesian Network, and each feature set instance is associated with precisely one outcome category.

Since the binning algorithm is dynamic, the bin sizes and ranges vary from dataset to dataset. As a result, the categories in each user-based dataset are different from one another, reflecting the different number of retweets that each author’s Tweets are likely to receive.

### 6.2.2 Varying Bin Sizes

Number of bins: explain how accuracy worsens as bin number increases.

## 6.3 Techniques Used

For the work presented in this chapter, various techniques and approaches were used for handling the data and for generating the scores. The most notable are now assessed in this section.

### 6.3.1 Machine Learning

Explain about Machine Learning, its uses, techniques and how this is useful.

Talk about how it was used in previous chapter, but that more in depth here. Bayesian Network

#### Classification Performance

Compared several types of classifiers (show table comparing accuracy, etc. of different types)

Explain that Bayesian Network is best (quick, accurate)

#### Training Results

Playing with Weka to improve the prediction performance (i.e. different number of bins, different features)

### 6.3.2 Crowdsourcing

Crowdsourcing is a technique often employed by companies and researchers to collect large amounts of data from a (often large) distribution of people or devices. Examples of this in practice include the use of surveys and user-contributed reviews (such as in Google Maps).

In this paper we use Amazon’s Mechanical Turk service as a method for crowdsourcing interestingness validations from many people, who are known as Mechanical Turk Workers (MTWs)

## 6.4 Experimentation

In this section we describe the stages involved in making interestingness inferences of Tweets based on their individual retweet behaviour. We start by discussing how the data was collected and how the information was then processed. We go on to illustrate how we employed our Bayesian Networks and the features we extracted for training them in order to make the retweet predictions. We finish by discussing how we then infer the Tweet interestingness and how this was validated.

### 6.4.1 Data Collection

In March 2013, a random walk was carried out through the Twitter social graph using Twitter’s REST API, originating with one of this paper’s author’s Twitter account. Each step of the walk involved focusing on one Twitter user, collecting information on that user, and then selecting a random follower of that user. This follower then became the focus of the next step.

At each step, a set of the most recent Tweets from the current user were collected. The number of Tweets returned by the Twitter API differed from user to user, based on their recent Tweet-posting frequency, though usually a few hundred Tweets were

yielded. In addition to their Tweets, we also collected information on the user itself and on a sample subset (up to 100, if they exist) of its friends and followers. We used a sample instead of collecting information on *all* of a user’s friends and followers so that we could maximise the efficiency of our method in terms of data collection. It also meant we had a snapshot of an additional 200 users in the author’s local network to give the classifier a notion of the activity of this local network both upstream and downstream from the author.

This gave a dataset containing around 241,000 Tweets from 370 unique Twitter users and, of those Tweets, around 90,000 were cases where the retweet volume was greater than zero. The dataset was split into two datasets: a training set (90%) and the testing set (10%). The original dataset was split in such a way as to allow all Tweets belonging to one particular user to exist in only one of the two smaller datasets. The training set was used to train a model, as described below, and was then discarded from the rest of the experimentation.

### 6.4.2 Data Corpora

As discussed in the methodology section, we are interested in producing *two* predictions for each Tweet; one when compared to the global set, and another from comparisons to the rest of that user’s Tweets. Therefore, a new dataset was formed for each user in the testing dataset which contained the Tweets only from that particular user. For the remainder of this section, the testing dataset is the *global* corpus of Tweets, and each of the individual user datasets are known as *user* corpora.

The Bayesian Network was chosen as our classifying algorithm to produce the predictions as it was suitable for the data types of the Tweet and user features and was also found to perform efficiently in terms of precision and recall across the outcome predictions when carrying out test cross-validations against itself. The prediction precision weighted across the outcome categories was calculated to be around 70% in cross-validation tests on the training dataset features.

### 6.4.3 Features

To train the Bayesian Network model, a series of new features were harvested. Generally, each of these features fell into one of two categories; user features and tweet features.

The Tweet features follow the same ideas as the features used in the previous chapter: static, generally binary features that describe the structure of the Tweet. The user (or ‘network’) features are related more to the *network* to which the Tweet belongs.

Features were extracted from the Tweet and user data to train the Bayesian Networks. Each Tweet is represented by an *instance* of feature sets specific to that Tweet (and its author, if appropriate). In each instance the outcome feature is the retweet volume, which is categorised using the technique discussed later.

Since some features are static among Tweets from the same user (for example, a user’s follower count), a different feature set is used to train the user models and the global model, as described below.

#### Features for the global corpus model

A total of 31 features were used to train the model classifier from the global data corpus.

The features used to train the global model are outlined in Table 6.1. The network features apply to both the followers and friends retrieved for each author. For example, the first feature of this category, ‘max. follower count’, represents two features referring to the maximum follower count observed in the sample of the author’s followers and in the sample of the author’s friends.

While Tweet features are permanent after the Tweet has been created, the author and network features are dynamic in that the social graph is of a constantly altering structure as links are formed and broken between users. However, in this work, we assume

Feature category	Feature	Feature data type
Tweet (‘genome’)	mention	{True, False}
	Tweet length	real (numeric)
	url	{True, False}
	hashtag	{True, False}
	positive emoticon	{True, False}
	negative emoticon	{True, False}
	exclamation mark	{True, False}
	question mark	{True, False}
	starts with ‘RT’	{True, False}
	is an @-reply	{True, False}
Author	follower count	real (numeric)
	friend count	real (numeric)
	verified account	{True, False}
	status count	real (numeric)
	listed count	real (numeric)
Network (‘environment’)	max. follower count	real (numeric)
	min. follower count	real (numeric)
	avg. follower count	real (numeric)
	max. friend count	real (numeric)
	min. friend count	real (numeric)
	avg. friend count	real (numeric)
	avg. status count	real (numeric)
	proportion verified	real (numeric)

**Table 6.1: Features used to train the model from the global data corpus**

that the changes to these features are not significant over the Tweets collected for each user and try to minimise this by only considering recent Tweets of each user.

### Features for individual user models

The 10 Tweet features were used for training each user-based Bayesian Network to reduce redundancy in the model since the author and network features would be the



same across all of the Tweets from one user.

#### 6.4.4 Making predictions and inferring Tweet interestingness

The global and user models were then trained using the features to produce one global model and one user model for each unique author in the testing dataset, as described above.

To produce the global predictions,  $T_{PG}$ , each Tweet in the testing database was evaluated against the global model, which output a predicted value for the bin that Tweet is predicted to belong to.

The user predictions,  $T_{PU}$  were calculated in the same way: each Tweet in the testing dataset was evaluated against the model based on the Tweet’s author.

Each Tweet then each had two retweet volume predictions. The upper bound of the binned category assigned to each prediction was compared to the *observed* retweet count to produce the two interestingness scores,  $TS_G$  and  $TS_U$ , as discussed earlier.

#### 6.4.5 Planning the Result Validation

To assist in validating the interestingness scores produced for the Tweets, tests were required in which humans would be responsible for also assessing how interesting they feel the Tweets were. The predicted interesting Tweets could then be compared to the Tweets chosen as interesting through human judgement.

For this stage of the experimentation, certain Tweets (and users) were removed from the dataset to be tested by the humans. Since the Tweet data was collected using a random walk through Twitter, there was no governance over the content of the text in this data. Therefore, users who frequently used offensive or non-English (since the Tweets were to be validated by English-speaking individuals) language had their Tweets removed from the test set. In addition, individual ‘@-replies’ (Tweets starting with the ‘@’ symbol) were also stripped from the test set since these usually denote

conversations on Twitter between two or more users and are unlikely to convey any interest to outsiders.

### Validating the Results

In this context, MTWs have no connection to the Tweets in the dataset, and are random users of Amazon’s Mechanical Turk who decide to take the job on. Collecting the validation data in this way, and instead of using a small set of people, means that we can collect a diverse opinion on Tweet interestingness from people whom the Tweets will have varying relevance to. This helps reinforce the notion of *global* interest levels, where Tweets can be interesting to a large number of users.

The human validations were carried out such that the MTWs were presented with questions, each of which consisting of five different Tweets from one specific author. Each question asked that the MTW select the Tweets they consider to be the most interesting of the group and that they must select at least one. MTWs were paid \$0.05 for each question they answered and there were 91 unique MTWs in total who responded to the questions.


The tests were under the conditions of a randomised controlled trial, such that each Tweet was assessed in three different contexts (i.e. each Tweet would appear in three different questions alongside four randomly chosen other Tweets) and that each question would be responded to by three different MTWs.

To represent the validation for this ongoing work, 750 Tweets were selected, at random, from the stripped testing dataset and were organised by user to ensure that Tweets only appeared in questions alongside other Tweets from the same author. Since each Tweet was required to appear in three different questions and each question consisted of 5 unique Tweets, this resulted in a total of 450 unique questions, each of which was answered by three different MTWs.

### 6.4.6 Validation Results

From the 450 questions asked of the MTWs, 325 questions had responses where the response was chosen with a confidence of two-thirds or greater. Since MTWs had the opportunity to choose more than one Tweet to be the most interesting, 349 Tweets (the subset of all tested Tweets), denoted as  $T$ , were selected as sufficiently interesting. Tweets selected from individual questions by only one of the assessing MTWs for that question were discarded.

#### General Performance



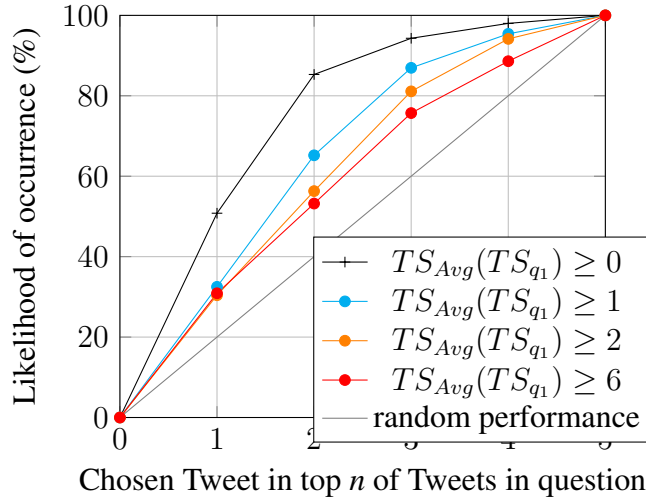
**Figure 6.1: Proportionate frequency distribution of global scores in the entire test set compared to only those in the subset  $T$ .**

In the subset  $T$ , our method predicted 140 Tweets to have  $TS_G(t) > 1$  and 80 Tweets where  $TS_U(t) > 1$ . These results were validated by the MTWs in that they agreed that all Tweets  $t \in T$  where  $TS_{Avg}(t) > 1$  were interesting in 60% of cases, with the performance of the global scores,  $TS_G(t)$ , achieving 65% precision. The user scores,  $TS_U(t)$ , were less accurate and yielded a precision of around 55%.

In addition, we show how the proportionate frequency of Tweets with higher values of

$TS_G(t)$  is greater in the subset  $T$  than in the entire set tested with the MTWs (Figure 6.1). This also demonstrates that Tweets with low global scores ( $0 \leq TS_G(t) < 1$ ) are more frequent in the entire test set than in the subset  $T$ .

### Per-Question Performance



**Figure 6.2:** The probability of the MTW's chosen Tweet's  $TS_{Avg}(t)$  being in the top  $n$  of Tweets for that question while varying the minimum allowed maximum Tweet score of the question..

Analyses were also conducted on the performance of the predictions on a per-question basis. Hereafter, a question,  $q \in Q$ , is defined as a set of Tweets where;

$$q = t_1^q, t_2^q, t_3^q, t_4^q, t_5^q$$


and

$$|q| = 5$$

For conducting these question-based analyses, each question's five Tweets were ranked in order of ascending mean interestingness score such that  $TS_{Avg}(t_1^q) \geq TS_{Avg}(t_2^q) \geq \dots \geq TS_{Avg}(t_5^q)$ . We then calculated the number of times the MTWs chose a Tweet that appeared in the top  $n$  of the ranked list of Tweets, as shown in Figure 6.2.

In this figure, we vary the minimum allowed value of  $TS_{Avg}(t_1^q)$  (the highest Tweet score in question  $q$ ) to show how detecting more interesting Tweets was more accurate when the range of scores in each question is more disparate. We show, for cases in which  $TS_{Avg}(q_1) \geq 1$ , that the likelihood of the MTWs choosing one of our two most highly ranked Tweets of the question using the interestingness predictions is around 66% and the chance that they choose one of the top three ranked Tweets is 87%.

### Probability of Selection



**Figure 6.3: The probability that Tweet  $t$  is chosen given that  $TS_G(t)$  is greater than a given value,  $x$ .**

Finally, we show that the probability of MTWs deciding a Tweet  $t$  is interesting becomes higher as the value of  $TS_G(t)$  increases.

In Figure 6.3, although we exclude cases of Tweets with  $TS_G(t) > 4$  to reduce noise (due to fewer samples), a significant increase in probability is observed, particularly in the interval 0-1 representing the range of Tweets that are uninteresting (fewer observed retweets than predicted) to ‘as expected’ (observed retweet volume is equal to predicted). It is also illustrated that, from this initial work, Tweets with a predicted interestingness score of 3 or more are not significantly different from one another in terms of their ‘real’, human-judged, interestingness. However, more work will be carried out towards this in the future so that more accurate research can be done on Tweets

with scores of greater than 3 in this context.

## 6.5 Further Analyses

The interestingness scores have been validated in terms of there being recognition that they can signify interesting Tweets. This section will now continue onto some deeper analyses of the results in order to show *how* it is able to work.

### 6.5.1 Discerning Interesting Information from Noise

In this subsection, the human-selected interestingness selection will be assessed. Of particular interest is the likelihood of humans agreeing on an interesting piece of information and the properties of the Tweet scores in questions when agreements *are* made.

For this, analyses were made into the *disparity* of scores for Tweets. That is, the range of scores of Tweets in a particular question and the effect this has on human decision in deducing interesting information.

Num. confident answers in $q$	min. $d^{Avg}(q)$	max. $d^{Avg}(q)$	avg. $d^{Avg}(q)$
0	0	846	<b>17.6</b>
> 0	0	1445	32.1
1	0	1445	<b>34.3</b>
> 1	0	4	0.647
> 2	0	0.55	0.204

**Table 6.2: Absolute  $TS_{Avg}(t)$  disparity of questions with varying number of confident answers made. Entries in bold are used to highlight interesting values..**

The absolute Tweet score disparity for a question,  $q$ , is defined as  $d(q)$ . For example, for a disparity of average scores, this is calculated thus;

$$d^{Avg}(q) = \max(TS_{Avg}(q)) - \min(TS_{Avg}(q))$$


Table 6.2 shows how the values for  $TS(q)$  vary for questions with differing numbers of confident answers. A confident answer, as mentioned previously, is one where at least two MTWs have agreed on an interesting Tweet.

The data shows that the average  $TS(q)$  is around double in cases where a question is answered with precisely one confident answers than in cases where there are no confident answers made. This indicates that a wider scale of interestingness in a question is useful for humans for picking out the content they'd prefer to read. If several pieces of content are more similarly interesting (or, as the case may be, uninteresting), then it becomes more difficult for an agreement to be made on which information is the *most* interesting.

In addition, the average score disparities in cases where multiple confident answers were selected are very low. This helps to reinforce the notion that pieces of information that are very similar in terms of interest level make it hard for users to decide on the *most* interesting. Indeed, in questions where this is the case, MTWs have selected, and agreed on, multiple Tweets.

To take this further, it is demonstrable that the probability of a confident selection being made for a particular question,  $q$ , increases as  $TS(q)$  also increases (Figure 6.4). Thus, this reinforces the notion that people find it easier to discern interesting information when compared to *un*-interesting information. This, too, is highlighted in Table 6.3, in which it is shown that amongst *all* questions (i.e. not only questions that have been confidently-answered) the score disparity is much smaller between Tweets that were selected than the score disparity for the entire question.

This is particularly the case in which there are a few Tweets which have similarly high scores amongst Tweets which collectively have *lower* scores. Therefore, selecting confidently from the few Tweets with the similar scores become difficult, but it is demonstrated that these at least are *more* interesting than the ones that weren't selected.



**Figure 6.4:** The probability of a confident selection being made for question  $q$  with varying  $d(q)$ .

	$TS_G(t)$	$TS_U(t)$	$TS_{Avg}(t)$
$TS_d(q)$	62.4	4.7	33.3
$TS_{d_C}(q)$	35.3	3.1	19.0
Ratio	57%	66%	58%

**Table 6.3:** Score disparity comparison between selected Tweets of question  $q$  and *all* Tweets in  $q$  when using the three different Tweet scores as metrics.

For example, this data shows that, on average, the global score disparity for selected Tweets of a question,  $q$ , was only around 57% that of the entire disparity of  $q$ .



---

## **Chapter 7**

# **Critical Assessment and Conclusions**

## **7.1 Critical Analysis of Results**

### **7.1.1 Analysis of Initial Research**

- What was the use of initial research?
- Are the results sensible?
- How have the results shaped the further research?

### **7.1.2 Analysis of Final Results**

- Have methods been able to sensibly predict retweet volumes?
- Have methods sensibly inferred Tweet interestingness?
- What might have worked better?
- Which parts were useless?
- Which parts helped develop other areas of research which may have provided further avenues of research ideas?

## 7.2 Further and Future Work

How can this research be taken further in the future?

- Use previous results to predict how far a tweet is likely to be retweeted (for advertising purposes)
- Useful for detecting the kind of messages that are likely to travel further
- As well as providing an interest level, the systems also predict sensible estimations on retweet volumes.
- Perhaps useful for measuring the spread of rumours.

## 7.3 Conclusions

### 7.3.1 Summary

Summarise events and processes covered, reiterate what the point of the work was and how each part of the work covered relates to that.

### 7.3.2 Contributions

Restate the original contributions (from Introduction section). Explain the ways in which the work done relates to the projected contributions, that it is novel and useful.

# Appendices

Source code, further diagrams, ideas, etc.

## Bibliography

- [1] Stuart M. Allen, Gualtiero Colombo, and Roger M. Whitaker. Uttering: social micro-blogging without the internet. In *Proceedings of the Second International Workshop on Mobile Opportunistic Networking*, MobiOpp '10, pages 58–64, New York, NY, USA, 2010. ACM.
- [2] Paavo Arvola, Jaana Kekäläinen, and Marko Junkkari. Expected reading effort in focused retrieval evaluation. *Information Retrieval*, 13:460–484, 2010. 10.1007/s10791-010-9133-9.
- [3] C Castillo, M Mendoza, and B Poblete. Information Credibility on Twitter. In *Proceedings of the 20th International Conference on World Wide Web*, WWW '11, pages 675 – 684. ACM, 2011.
- [4] Meeyoung Cha, Hamed Haddadi, Fabricio Benevenuto, and Krishna P. Gummadi. Measuring User Influence in Twitter: The Million Follower Fallacy. In *Fourth International AAAI Conference on Weblogs and Social Media*, May 2010.
- [5] MJ Chorley, GB Colombo, SM Allen, and RM Whitaker. Better the Tweeter You Know: Social Signals on Twitter. In *International Conference on Social Computing*, pages 277 – 282. ASE/IEEE, 2012.
- [6] A. Clauset, C. Rohilla Shalizi, and M. E. J. Newman. Power-law distributions in empirical data. *ArXiv e-prints*, June 2007.
- [7] R Dawkins. *The Selfish Gene*. Oxford University Press, 1976.
- [8] Wojciech Galuba, Karl Aberer, Dipanjan Chakraborty, Zoran Despotovic, and Wolfgang Kellerer. Outtweeting the twitterers - predicting information cascades in microblogs. In *Proceedings of the 3rd conference on Online social networks*, WOSN'10, pages 3–3, Berkeley, CA, USA, 2010. USENIX Association.

- [9] Oliver Hein, Michael Schwind, and Wolfgang König. Scale-free networks. *WIRTSCHAFTSINFORMATIK*, 48:267–275, 2006. 10.1007/s11576-006-0058-2.
- [10] L Hong, O Dan, and B D Davison. Predicting Popular Messages in Twitter. In *Proceedings of the 20th International Conference Companion on World Wide Web*, WWW '11, pages 57 – 58. ACM, 2011.
- [11] A Java, X Song, T Finin, and B Tseng. Why We Twitter: Understanding Microblogging Usage and Communities. In *Proceedings of the 9th WebKDD and 1st SNA-KDD 2007 workshop on Web mining and social network analysis*, WebKDD/SNA-KDD '07, pages 56 – 65. ACM, 2007.
- [12] H Kwak, C Lee, H Park, and S Moon. What is Twitter, a Social Network or a News Media? In *Proceedings of the 19th international conference on World wide web*, WWW '10, pages 591 – 600. ACM, 2010.
- [13] N Naveed, T Gottron, J Kunegis, and A C Alhadi. Bad News Travel Fast: A Content-based Analysis of Interestingness on Twitter. In *In: Proceedings of the ACM WebSci'11*, pages 1 – 7. ACM, 2011.
- [14] H-K Peng, J Zhu, D Piao, R Yan, and Y Zhang. Retweet Modeling Using Conditional Random Fields. In *Data Mining Workshops (ICDMW), 2011 IEEE 11th International Conference on*, pages 336 – 343. IEEE, 2011.
- [15] Bongwon Suh, Lichan Hong, P. Pirolli, and E.H. Chi. Want to be retweeted? large scale analytics on factors impacting retweet in twitter network. In *Social Computing (SocialCom), 2010 IEEE Second International Conference on*, pages 177 –184, aug. 2010.
- [16] I Uysal and W B Croft. User Oriented Tweet Ranking: a Filtering Approach to Microblogs. In *Proceedings of the 20th ACM international conference on Information and knowledge management*, CIKM '11, pages 2261 – 2264. ACM, 2011.
- [17] W Webberley, S Allen, and R Whitaker. Retweeting: A study of Message-Forwarding in Twitter. In *Workshop on Mobile and Online Social Networks (MOSN)*, pages 13 – 18. IEEE, 2011.

- 
- [18] Y Xu. Relevance Judgment in Epistemic and Hedonic Information Searches. *Journal of the American Society for Information Science and Technology*, 58(2):179–189, 2007.
  - [19] Dejin Zhao and Mary Beth Rosson. How and why people twitter: the role that micro-blogging plays in informal communication at work. In *Proceedings of the ACM 2009 international conference on Supporting group work*, GROUP '09, pages 243–252, New York, NY, USA, 2009. ACM.
  - [20] J Zhu, F Xiong, D Piao, Y Liu, and Y Zhang. Statistically Modeling the Effectiveness of Disaster Information in Social Media. In *Global Humanitarian Technology Conference (GHTC)*, pages 431 – 436. IEEE, 2011.